

1 Ideas/Issues

Even if I train the observation model and the controls at the same time successfully, how do I motivate the controls to do what I want?

Use the built in reward (+1 for every time frame where the pole is "up")?
Add an extra term to loss that depends on both time and controls?

Devise a better loss function. The current loss function is $\mathcal{L}(\phi) = \sum_{i=1}^N \log p_{SS}(z_{1:T_i} | \Theta_{1:T_i})$

Maybe use unscented Kalman Filter??

Multiply the loss term by a factor of $1/t$ or $1/t^2$. Something along those lines to get implicit reliance on time component. Suggests that a system that learns to stay upright longer has less error/smaller gradient than a system that falls over quickly.

Run x trials with the current model, train the model on those trials with the loss multiplier. Repeat with the updated parameters.

Need a separate optimization problem for learning the control? Maybe learn mapping from observations to control directly with nn.
Design a policy function/mapping from observations to controls

1.1 Training

The considerations for training are:

- What/when I'm training (observation model, control,
- what the goal of that training is(learn to control the system, or just explore, i.e. learn to predict the next state)
- How to get training samples(online/offline)
- LSTM fixed/variable length
- Inputs to LSTM(physical specifications, previous outputs, etc.)

1.2 Thoughts

- Train transition model with LQR, hope that the model learns how to linearize given any particular observation. Then separately train the learned transition/control models to put the pole into any place of our choosing. Loss function: euclidean distance from our desired state? Only works if we have knowledge of what the states mean

- Suppose that I include the previous state as input to the LSTM, the system becomes auto-regressive, do the assumptions leading to the kalman filter, and the log-likelihood function still hold?

2 Assumptions

We have the linear dynamical system:

$$\begin{aligned} l_{t+1} &= A_t l_t + B_t u_t + g_t \epsilon_t & \epsilon_t &\sim \mathcal{N}(0, 1) \\ Z_t &= C_t l_t + D_t u_t + \sigma_t \varepsilon_t & \varepsilon_t &\sim \mathcal{N}(0, 1) \end{aligned}$$

where $A \in \mathbb{R}^{m \times n}$, $B \in \mathbb{R}^{m \times r}$, $C \in \mathbb{R}^{z \times m}$, $D \in \mathbb{R}^{z \times r}$, $l \in \mathbb{R}^m$, $u \in \mathbb{R}^r$, $Z \in \mathbb{R}^z$, $\epsilon \in \mathbb{R}^m$, $\varepsilon \in \mathbb{R}^z$ However, it simplifies if, as in our situation, $C = I$, and $D = 0$.

3 Potential mistakes

In computing the likelihood: $\mathcal{N}(z_t | \mu_t, \Sigma_t)$ Compare the paper's calculation:

$$\begin{aligned} \mu_1 &= a_1^T \mu_0 \\ \Sigma_1 &= a_1^T \Sigma_0 a_1 + \sigma_1^2 \\ \mu_t &= a_t^T F_t f_{t-1} \\ \Sigma_t &= a_t^T (F_t S_t F_t^T + g_t g_t^T) a_t + \sigma_t^2 \end{aligned}$$

with my control version:

```
mu_1 = tf.matmul(trans(self.C)[0], self.mu_0)
mu = tf.matmul(C, tf.add(tf.matmul(A, l_filtered), tf.matmul(B, u)))

Sigma_1 = tf.matmul(tf.matmul(trans(self.C)[0],
    tf.linalg.diag(tf.squeeze(self.Sigma_0))), trans(self.C)[0],
    transpose_b=True) + tf.square(trans(self.sigma)[0])

temp = tf.matmul(tf.matmul(A, P_filtered), A, transpose_b=True) +
    tf.matmul(g, g, transpose_b=True)
Sigma = tf.matmul(tf.matmul(C, temp), C, transpose_b=True) +
    tf.square(sigma)
```
